

Job Title: Senior Machine Learning Engineer / Data Scientist

Company: Horizon Data Intelligence

Experience: 3-6 Years

Location: Remote

Overview:

We are seeking an experienced ML Engineer to develop, train, and deploy predictive models and deep learning pipelines. You will work on tabular data, NLP transformers, and model serving systems.

Key Responsibilities:

- Design end-to-end Machine Learning pipelines from data preprocessing to model deployment.
- Train and fine-tune models using PyTorch, TensorFlow, Scikit-Learn, and Hugging Face Transformers.
- Build scalable feature engineering workflows with Pandas, NumPy, and Apache Spark.
- Implement model evaluation metrics (AUC-ROC, F1-Score, RMSE) and hyperparameter tuning.
- Deploy models as REST microservices with FastAPI, Docker, and MLflow / Triton Inference Server.
- Collaborate with software engineers to integrate models into production environments.

Requirements:

- Core Languages: Python, SQL
- ML/DL Frameworks: PyTorch, Scikit-Learn, Hugging Face, XGBoost, LightGBM
- Data Tools: Pandas, NumPy, Apache Spark
- MLOps & Deployment: Docker, FastAPI, MLflow, AWS SageMaker
- Strong understanding of statistical modeling, linear algebra, and data validation.